Table of Contents

1. Chapter 1: Megabrains 101: Generative AI and LLMs Unboxed
    1. Understanding Gen AI
    1. Understanding LLMs
    1. Understanding the Role of Hugging Face in the LLM World
    1. Ecosystem Around LLMs
    1. Challenges and Considerations in Using LLMs
    1. Setting the Stage: LLMs in Java Development
    1. Conclusion

1. Chapter 2: First Contact: “Hello, LLM” with Spring Boot
    1. Introduction
    1. What Is OpenAI?
    1. Creating Your First Spring Boot LLM Application
    1. Exploring the OpenAI Chat API
    1. Integrating OpenAI API with Spring Boot
    1. Conclusion

1. Chapter 3: The Transformer Saga—From Attention to Fine-Tuning
    1. Introduction
    1. Importance of LLMs
    1. How Large Language Models Work
    1. Transformer
    1. Evolution of Transformers
    1. Fine-Tuning LLMs
        1. Supervised Fine-Tuning
        1. Dialogue and Instruction Tuning
        1. Safety Tuning
        1. Reinforcement Learning from Human Feedback
        1. Parameter-Efficient Fine-Tuning
    1. Inference Optimization Techniques
        1. Quantization
        1. Distillation
        1. Output-Preserving Optimizations

        1. Speculative Decoding
        1. Batching and Parallelization
    1. Conclusion

1. Chapter 4: Bring Your Own Model: Self-Hosting with Ollama
    1. Introduction
    1. Self-Hosting LLMs
    1. Ollama
    1. Installing Ollama
    1. Running a Locally Hosted LLM
    1. Exploring the APIs Exposed by Ollama
    1. Integrating Ollama with Your Existing Application
    1. Conclusion

1. Chapter 5: Power Tools: LangChain4j Quick-Start
    1. Introduction
    1. What Is LangChain4j?
    1. Installation and Setup of LangChain4j
    1. Initial Project Setup: Creating Your First LangChain4j Application
    1. Conclusion

1. Chapter 6: Integrating LLMs with Java Applications
    1. Introduction
    1. Prompt Engineering
    1. Context into Conversation Memory
    1. Switching Between LLM Providers
    1. Handling LLM Outputs
    1. Security Considerations
    1. Conclusion

1. Chapter 7: From Chatty to Clever: Retrieval-Augmented Generation
    1. Introduction
    1. Contextuality in the LLM Ecosystem
    1. RAG Paradigm
    1. Embeddings
    1. Vector Store
    1. RAG in Action
    1. Conclusion

1. Chapter 8: Spring AI Ninja Moves

    1. Introduction
    1. What Is Spring AI?
    1. Features of Spring AI
    1. Boot Up an Application Using Spring AI
    1. Using OpenAI API Through Ollama API in Spring AI
    1. Building a RAG-Based Application
        1. Retrieval
    1. Conclusion

1. Chapter 9: Prompt Alchemy: Patterns That Make Models Look Smarter
    1. Introduction
    1. What Is a Prompt?
    1. What Is Prompt Engineering?
    1. Prompt Engineering Patterns
        1. Creating Effective Prompts
        1. Simple Prompting Techniques
        1. Advanced Prompting Techniques
        1. Microsoft's Framework for Prompt Creation and Optimization
        1. Best Practices for Prompt Engineering
    1. Using Prompt Engineering in Spring AI
        1. Temperature
        1. Output Length (MaxTokens)
        1. Structured Response Format
        1. Model-Specific Options
    1. Prompt Engineering Techniques
        1. Zero-Shot Prompting
        1. One-Shot and Few-Shot Prompting
        1. System, Contextual, and Role Prompting
        1. Step-Back Prompting
        1. Chain of Thought
        1. Self-Consistency
        1. Tree of Thoughts
        1. Automatic Prompt Engineering
        1. Code Prompting
    1. Notes
    1. Conclusion

1. Chapter 10: Swiss-Army LLMs: Tool Calls in Spring AI
    1. Introduction
    1. What Are Tools?

1. Understanding the Flow of Control with Tools
1. Understanding Tool Calling with Spring AI
1. Standard Practices in Tool Calling
1. Starting with a Simple Use of Tools
    1. Declarative Specification: @Tool
1. Exploring an Advanced Use Case of Tools
1. Conclusion

1. Chapter 11: Agents Assemble! Building Autonomous Workflows
    1. Introduction
    1. Agentic AI
        1. Key Characteristics of an Agent
        1. Types of Agents
        1. Applications of Agents
    1. Defining Concepts
    1. Components of Agentic Systems
        1. Models
        1. Tools
        1. Knowledge and Memory
        1. Guardrails
        1. Orchestration
    1. Agentic Workflows and Patterns
        1. Chain Workflow
        1. When to Use Chain Workflow
        1. What's Happening in the Workflow?
        1. Parallelization Workflow
        1. What's Happening in the Process?
        1. Routing Workflow
        1. When to Use
        1. What's Happening in the Process?
        1. Key Insights:
    1. Agentic Workers
        1. Orchestrator-Workers
        1. When to Use
        1. OrchestratorWorkersWorkflow Class
        1. What's Happening in the Process?
        1. Key Insights
        1. Evaluator-Optimizer
        1. When to Use
        1. EvaluatorOptimizerWorkflow Class

        1. What's Happening in the Process?
        1. Key Insights
    1. Agent Lifecycle in Spring AI
    1. Spring AI's Implementation Advantages
        1. Model Portability
        1. Structured Output
        1. Consistent API
    1. Best Practices and Recommendations
        1. Start Simple
        1. Design for Reliability
        1. Consider Trade-Offs
    1. Advisors API in Spring AI
    1. Conclusion

1. Chapter 12: Quarkus + LangChain4j: Lightning-Fast Gen AI
    1. Introduction
    1. Integration of LangChain4j in Quarkus
    1. Prompt Engineering in Quarkus
        1. Crafting Prompts with Annotations
        1. Exposing an AI-Powered REST Endpoint
    1. Using Tools in Quarkus
    1. Using ChatMemory in Quarkus
        1. Defining a ChatMemoryProvider
        1. Enabling Memory in Your AI Service
    1. Using the RAG Pattern in Quarkus
        1. Ingest Documents into Your Redis Vector Store
    1. Conclusion

1. Chapter 13: Jlama and Friends: Hosting Models the Java Way
    1. Introduction
    1. What Is Jlama?
    1. Jlama: Quick Start
    1. Use Jlama to Host an LLM
    1. Quantize a Model Using Jlama
    1. Use Jlama with Quarkus to Build a Gen AI Application
    1. Conclusion

1. Chapter 14: Seeing Is Believing: Multimodal LLMs and Image Hacking
    1. Introduction
    1. Tooling in Java for Image Processing

    1. Multimodal LLMs
    1. Building an Application to Process Flight Tickets
    1. Generating Images from User Request
    1. Other multimodal formats
    1. Conclusion

1. Chapter 15: Does It Even Work? Testing and Evaluating LLM Apps
    1. Introduction
    1. Evaluation of LLM Applications
    1. Tools Available for Evaluation and Testing
    1. Testing and Evaluating Using Spring AI
    1. Conclusion

1. Chapter 16: Cloud Power-Ups: Bedrock, Vertex, and Azure OpenAI
    1. Introduction
    1. State of Gen AI in the Cloud
    1. Generative AI in AWS: Amazon Bedrock
    1. Accessing a Foundational Model in Bedrock Playground
    1. Building a Sample Project Using Amazon Bedrock SDK
        1. Using AWS SDK to Connect with Amazon Bedrock
    1. Extending and Integrating Your Application with Guardrails
        1. Amazon Bedrock Guardrails: Enabling Responsible AI at Scale
    1. Extending and Integrating RAG into Your Application
    1. Reengineering Your Application Using Spring AI Integration with Amazon Bedrock
    1. Generative AI in GCP
    1. Generative AI in Azure
    1. Conclusion

1. Chapter 17: Talking in Protocols: The MCP Revolution
    1. Introduction
    1. Protocols in the LLM Ecosystem
        1. Model Context Protocol
        1. Agent Communication Protocol
        1. Agent-to-Agent Protocol
        1. Agent Network Protocol
    1. What Is MCP Protocol?
    1. Role of MCP
    1. MCP Using Spring AI
        1. MCP Client

        1. MCP Server
        1. Spring AI MCP Integration
        1. Capabilities of the MCP Server in Spring AI
    1. Creating Spring AI-Based MCP Server
    1. Creating a Spring AI-Based MCP Client
    1. Integrating MCP with Quarkus
    1. Conclusion

1. Chapter 18: Can You See Me Now? Observability for LLM Pipelines
    1. Introduction
    1. Why Observability Matters in LLM Evaluation
    1. Core Observability Signals
    1. Best Practices
    1. Integrating Observability in LLMs Through Spring AI
        1. Observing ChatClient Requests
        1. Observing Chat Model Requests
    1. Conclusion

1. Chapter 19: Native-Speed Machine Learning in Java: DJL, ONNX, and JNI
    1. Introduction
    1. Using Hugging Face APIs in Java
        1. Accessing Hugging Face Models via an API
        1. Creating Java Wrappers for Inference
        1. Integration via Frameworks
    1. Deep Java Library (DJL)
    1. Hosting and Inference Using DJL
    1. Understanding the Linnerud Dataset
        1. Dataset Structure
        1. Use in Machine Learning
        1. Why Use This Toy Dataset?
    1. Building an ONNX Regression Model with the Linnerud Dataset
        1. Step-by-Step: Training and Exporting the Model
        1. Why This Workflow Works
    1. Integrating LLMs Using Java Native Interface
        1. Why Use JNI for LLM Integration?
        1. Sample JNI Integration with a Native LLM
        1. Behind the Scenes
        1. When to Use JNI for LLMs
        1. Considerations
    1. Running Inference in Java with ONNX Runtime

        1. Why Use ONNX Runtime in Java?
        1. Getting Started with ONNX Runtime for Java
        1. Key Concepts
        1. Use Cases
    1. Model Training and Fine-Tuning Using DL4J
        1. Key Components of the DL4J Suite
        1. Training and Fine-Tuning Models
        1. Why Use DL4J?
        1. Open Source and Community
    1. Natural Language Processing Using Apache OpenNLP
        1. Getting Started
        1. Why Use OpenNLP?
    1. Conclusion

1. Chapter 20: Architectures of Tomorrow: From Monoliths to Modular Minds
    1. Introduction
    1. Architectural and Scaling Trends in Generative AI: From Monoliths to Modular Intelligence
        1. Smarter Scaling: From Large Models to Efficient Architectures
        1. RAG: Memory-Enhanced Systems for Real-Time Reasoning
        1. Agentic AI: Modular Orchestration Through Intelligent Agents
        1. Protocol-Driven Interoperability: MCP, ACP, A2A, and ANP
        1. Performance-Aware Design: Beyond the Model
        1. A New Architecture for Scalable Intelligence
    1. Ecosystem Showdown: Open-Source vs. Cloud-Hosted Models
    1. Navigating the Challenges of Generative AI: What Developers Must Know and Do
        1. Prompt Injection, Jailbreaks, and Model Exploits
        1. Ethical Use and Bias Mitigation
        1. Data Privacy and Leakage
        1. Governance and Compliance
        1. Energy Use and Environmental Cost
        1. Model Hallucinations and Factual Integrity
        1. Staying Current: The Developer's Playbook
    1. The Evolving Pathway to LLM Adoption: From Prompts to Agentic AI
        1. Phase 1: Prompt Engineering and Model Trials
        1. Phase 2: Fine-Tuning and Domain Adaptation
        1. Phase 3: Model Training from Scratch (Advanced Tier)
        1. The Rise of Agentic AI
        1. Managing the LLM Development Lifecycle

1. LLMOps: Operationalizing the Lifecycle of Generative AI Systems
    1. Core Components of LLMOps
    1. LLMOps in Practice: From Playground to Production
    1. The Road Ahead
1. Bridging the Gap: Migrating from DevOps to LLMOps Thinking
    1. From Deterministic Code to Probabilistic Outputs
    1. From CI/CD Pipelines to Continuous Evaluation
    1. From Logs and Metrics to Semantic Observability
    1. From Source Code to Prompt + Model as Code
    1. From Feature Testing to Use-Case Auditing
    1. From Code Deployments to Model Lifecycle Management
    1. From User Feedback to Human-in-the-Loop Learning
    1. Final Thought: It's Still About Culture
1. The Evolving Culture and Practices in Java's Machine Learning Ecosystem
    1. Java Meets Machine Learning: A Culture Shift in the Making
    1. DJL: Lowering the Barrier for Java Developers
    1. Eclipse Deeplearning4j: Full-Stack Deep Learning on the JVM
    1. Jlama and the Expanding Toolchain
    1. Where the Culture Is Heading
    1. A Rising Force in AI Development
1. Emerging Frontiers: Multimodality, Agents, and Domain Specialization
1. Regulations and Acts
    1. What This Means for Builders
1. Governance, Ethics, and Future Outlook
1. There's Just Too Much Happening. What Should I Do?
    1. Focus on Fundamentals, Not the FOMO
    1. You Already Have Transferable Knowledge
    1. Embrace the Culture of Learning
    1. Adaptation Is the Skill of the Decade
1. Conclusion

1. Index

Previous chapter

< Acknowledgments

Next chapter

About the Author >